

MRP - Week 8-9

This week we will be experimenting with using a variety of layers and α/β ratios

- The VGG-16 neural network allows us to extract the style and content of images separately
- The loss function is minimized during image synthesis and has 2 terms for content and style loss respectively, allowing us to control emphasis of one over the other.

Loss function

$$L_{total}(S, C, G) = \alpha L_{content}(C, G) + \beta L_{style}(S, G)$$

- The loss function needs to be fed a set of layers from the VGG-16 network
 - o $L_{content}$: takes any 1 layer of the network
 - o L_{style} : takes layers 'conv1 1', 'conv2 1', 'conv3 1', 'conv4 1' and 'conv5 1' (basically the 1st layer of each block of the VGG net to give a good variation of style details)
- The goal is to experiment generating an image using a variety of inputs layers for $L_{content}$ against a varying ratio of α/β



content

+



style

=

Generated
image

$$\alpha/\beta = 10^3$$






$$\alpha/\beta = 10^2$$

$$\alpha/\beta = 10^1$$





$$\alpha/\beta = 0.25$$

$$\alpha/\beta = 0.5$$





Con1_1



Con2_1



Con3_1



$$\alpha/\beta = 10^3$$

$$\alpha/\beta = 10^2$$

$$\alpha/\beta = 10^1$$


$$\alpha/\beta = 0.25$$

$$\alpha/\beta = 0.5$$

Con4_1



Con5_1



- The grid of experiments above show how the generated image varies with respect to the layer used for content representation
- All experiments were run without adding the total variation loss factor in order to better see the effect of change
- Note:
 - Higher layers capture high-level content features. Not exact pixel value
 - Higher layers are more concerned with the shapes and arrangement of objects in the image
 - Progressively, the higher layers show less resemblance to the original content image

- The following trends can easily be observed in the grid of images above:
 - o As α/β increases \rightarrow more emphasis on the content image
 - o As α/β decreases \rightarrow more emphasis on the style image
- This grid can be used to visually optimize the generated results by allowing us to select how much detail we wish to retain
- Content images containing subjects like faces or other fine details may require that we use layers from block 1, 2, or 3. No higher
- My results are comparable to the work of Leon Gatys' in his paper , "*A neural Algorithm of Artistic Style*" (see below)



The column headings read the α/β ratio values as 10^{-5} , 10^{-4} , 10^{-3} , and 10^{-2}

The row headings read the layers used as conv1_1, conv2_1, conv3_1, conv4_1, conv5_1. Basically the first layer of each block in the VGG network